

# Assignment 4: Data Wrangling (Fall 2025)

Isabella Yang

## OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on Data Wrangling

## Directions

1. Rename this file `<FirstLast>_A04_DataWrangling.Rmd` (replacing `<FirstLast>` with your first and last name).
2. Change “Student Name” on line 3 (above) with your name.
3. Work through the steps, **creating code and output** that fulfill each instruction.
4. Be sure to **answer the questions** in this assignment document.
5. When you have completed the assignment, **Knit** the text and code into a single PDF file.
6. Ensure that code in code chunks does not extend off the page in the PDF.

## Set up your session

- 1a. Load the `tidyverse`, `lubridate`, and `here` packages into your session.
  - 1b. Check your working directory.
  - 1c. Read in all four raw data files associated with the EPA Air dataset, being sure to set string columns to be read in a factors. See the README file for the EPA air datasets for more information (especially if you have not worked with air quality data previously).
2. Add the appropriate code to reveal the dimensions of the four datasets.

```
#1a Load packages
```

```
library(tidyverse)
```

```
library(lubridate)
```

```
library(here)
```

```
#1b Check working directory
```

```
here()
```

```
## [1] "/home/guest/EDE_Fall2025"
```

```
#1c Read all four raw data
```

```
folder <- here("Data", "Raw")
```

```
# List all four files that start with EPAair_
```

```
files <- dir(folder, pattern = "^EPAair_.*\\.csv$", full.names = TRUE)
```

```

# Read them in as a list of data frames
EPAair.data <- files %>%
  map(~ read.csv(.x, stringsAsFactors = TRUE))

# Name and read each file
EPAair_01 <- EPAair.data[[1]]
EPAair_02 <- EPAair.data[[2]]
EPAair_03 <- EPAair.data[[3]]
EPAair_04 <- EPAair.data[[4]]

#2 Check the dimensions of each dataset
dim(EPAair_01)

```

```
## [1] 9737 20
```

```
dim(EPAair_02)
```

```
## [1] 10592 20
```

```
dim(EPAair_03)
```

```
## [1] 8983 20
```

```
dim(EPAair_04)
```

```
## [1] 8581 20
```

All four datasets should have the same number of columns but unique record counts (rows). Do your datasets follow this pattern? Yes

## Wrangle individual datasets to create processed files.

3. Change the Date columns to be date objects.
4. Select the following columns: Date, DAILY\_AQI\_VALUE, Site.Name, AQS\_PARAMETER\_DESC, COUNTY, SITE\_LATITUDE, SITE\_LONGITUDE
5. For the PM2.5 datasets, fill all cells in AQS\_PARAMETER\_DESC with “PM2.5” (all cells in this column should be identical).
6. Save all four processed datasets in the Processed folder. Use the same file names as the raw files but replace “raw” with “processed”.

```

#3 Convert Date column to proper date format
library(lubridate)

fix_date <- function(x) {
  # x can be factor or character; return Date
  z <- as.character(x)
  # try multiple orders; handles with/without time

```

```

as.Date(parse_date_time(z,
                        orders = c("mdy", "ymd", "mdy HMS", "ymd HMS", "mdy HM", "ymd HM")),
      tz = "UTC")
}

# Re-read the four files (so we undo the NA overwrite) and parse Date
EPAair_01 <- read.csv(here::here("Data", "Raw", "EPAair_03_NC2018_raw.csv"), stringsAsFactors = TRUE)
EPAair_02 <- read.csv(here::here("Data", "Raw", "EPAair_03_NC2019_raw.csv"), stringsAsFactors = TRUE)
EPAair_03 <- read.csv(here::here("Data", "Raw", "EPAair_PM25_NC2018_raw.csv"), stringsAsFactors = TRUE)
EPAair_04 <- read.csv(here::here("Data", "Raw", "EPAair_PM25_NC2019_raw.csv"), stringsAsFactors = TRUE)

EPAair_01$Date <- fix_date(EPAair_01$Date)
EPAair_02$Date <- fix_date(EPAair_02$Date)
EPAair_03$Date <- fix_date(EPAair_03$Date)
EPAair_04$Date <- fix_date(EPAair_04$Date)

#4 Keep only the specified columns
EPAair_01 <- EPAair_01 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY, SITE_LATITUDE, SITE_LONGITUDE)
EPAair_02 <- EPAair_02 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY, SITE_LATITUDE, SITE_LONGITUDE)
EPAair_03 <- EPAair_03 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY, SITE_LATITUDE, SITE_LONGITUDE)
EPAair_04 <- EPAair_04 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY, SITE_LATITUDE, SITE_LONGITUDE)

#5 Update the AQS_PARAMETER_DESC for PM2.5 datasets
EPAair_03$AQS_PARAMETER_DESC <- "PM2.5"
EPAair_04$AQS_PARAMETER_DESC <- "PM2.5"

#6 Save processed datasets
write.csv(EPAair_01,
  here("Data", "Processed", "EPAair_03_NC2018_processed.csv"),
  row.names = FALSE)
write.csv(EPAair_02,
  here("Data", "Processed", "EPAair_03_NC2019_processed.csv"),
  row.names = FALSE)
write.csv(EPAair_03,
  here("Data", "Processed", "EPAair_PM25_NC2018_processed.csv"),
  row.names = FALSE)
write.csv(EPAair_04,
  here("Data", "Processed", "EPAair_PM25_NC2019_processed.csv"),
  row.names = FALSE)

```

## Combine datasets

7. Combine the four datasets with `rbind`. Make sure your column names are identical prior to running this code.
8. Wrangle your new dataset with a pipe function (`%>%`) so that it fills the following conditions:
  - Include only sites that the four data frames have in common:

“Linville Falls”, “Durham Armory”, “Leggett”, “Hattie Avenue”, “Clemmons Middle”, “Mendenhall School”, “Frying Pan Mountain”, “West Johnston Co.”, “Garinger High School”, “Castle Hayne”, “Pitt Agri. Center”, “Bryson City”, “Millbrook School”

(the function `intersect` can figure out common factor levels - but it will include sites with missing site information, which you don’t want...)

- Some sites have multiple measurements per day. Use the split-apply-combine strategy to generate daily means: group by date, site name, AQS parameter, and county. Take the mean of the AQI value, latitude, and longitude.
  - Add columns for “Month” and “Year” by parsing your “Date” column (hint: `lubridate` package)
  - Hint: the dimensions of this dataset should be 14,752 x 9.
9. Spread your datasets such that AQI values for ozone and PM2.5 are in separate columns. Each location on a specific date should now occupy only one row.
  10. Call up the dimensions of your new tidy dataset.
  11. Save your processed dataset with the following file name: “EPAair\_O3\_PM25\_NC1819\_Processed.csv”

```
#7 Combine the four datasets
# Check if the column names are identical
colnames(EPAair_01)
```

```
## [1] "Date"           "DAILY_AQI_VALUE"  "Site.Name"
## [4] "AQS_PARAMETER_DESC" "COUNTY"         "SITE_LATITUDE"
## [7] "SITE_LONGITUDE"
```

```
colnames(EPAair_02)
```

```
## [1] "Date"           "DAILY_AQI_VALUE"  "Site.Name"
## [4] "AQS_PARAMETER_DESC" "COUNTY"         "SITE_LATITUDE"
## [7] "SITE_LONGITUDE"
```

```
colnames(EPAair_03)
```

```
## [1] "Date"           "DAILY_AQI_VALUE"  "Site.Name"
## [4] "AQS_PARAMETER_DESC" "COUNTY"         "SITE_LATITUDE"
## [7] "SITE_LONGITUDE"
```

```
colnames(EPAair_04)
```

```
## [1] "Date"           "DAILY_AQI_VALUE"  "Site.Name"
## [4] "AQS_PARAMETER_DESC" "COUNTY"         "SITE_LATITUDE"
## [7] "SITE_LONGITUDE"
```

```
EPAair_all <- bind_rows(EPAair_01, EPAair_02, EPAair_03, EPAair_04)
```

```
dim(EPAair_all)
```

```
## [1] 37893      7
```

```
#8 Wrangle combined dataset
```

```
nrow(EPAair_01)
```

```
## [1] 9737
```

```
nrow(EPAair_02)
```

```
## [1] 10592
```

```
nrow(EPAair_03)
```

```
## [1] 8983
```

```
nrow(EPAair_04)
```

```
## [1] 8581
```

```
subset_test <- EPAair_all %>%  
  filter(Site.Name %in% c(  
    "Linville Falls", "Durham Armory", "Leggett", "Hattie Avenue",  
    "Clemmons Middle", "Mendenhall School", "Frying Pan Mountain",  
    "West Johnston Co.", "Garinger High School", "Castle Hayne",  
    "Pitt Agri. Center", "Bryson City", "Millbrook School"  
  ))
```

```
nrow(subset_test) # this return 16,510
```

```
## [1] 16510
```

```
# Split-apply-combine to daily means
```

```
EPAair_daily <- subset_test %>%  
  group_by(Date, Site.Name, AQS_PARAMETER_DESC, COUNTY) %>%  
  summarise(  
    DAILY_AQI_VALUE = mean(DAILY_AQI_VALUE, na.rm = TRUE),  
    SITE_LATITUDE = mean(SITE_LATITUDE, na.rm = TRUE),  
    SITE_LONGITUDE = mean(SITE_LONGITUDE, na.rm = TRUE),  
    .groups = "drop"  
  )
```

```
# Add Month/Year
```

```
EPAair_daily$Month <- lubridate::month(EPAair_daily$Date)
```

```
EPAair_daily$Year <- lubridate::year(EPAair_daily$Date)
```

```
EPAair_daily <- EPAair_daily %>%  
  select(Date, Month, Year, Site.Name, AQS_PARAMETER_DESC, COUNTY,  
    DAILY_AQI_VALUE, SITE_LATITUDE, SITE_LONGITUDE)
```

```
dim(EPAair_daily) # this should return 14,752 x 9
```

```
## [1] 14752      9
```

```
#9 Spreading datasets
library(tidyr)

EPAair_wide <- EPAair_daily %>%
  pivot_wider(
    names_from = AQS_PARAMETER_DESC,      # Ozone vs PM2.5
    values_from = DAILY_AQI_VALUE        # put AQI values in new columns
  )

#10
dim(EPAair_wide)
```

```
## [1] 8976    9
```

```
#11
write.csv(EPAair_wide,
  here("Data", "Processed", "EPAair_03_PM25_NC1819_Processed.csv"),
  row.names = FALSE)
```

## Generate summary tables

12. Use the split-apply-combine strategy to generate a summary data frame. Data should be grouped by site, month, and year. Generate the mean AQI values for ozone and PM2.5 for each group. Then, add a pipe to remove instances where mean **ozone** values are not available (use the function `drop_na` in your pipe). It's ok to have missing mean PM2.5 values in this result.

13. Call up the dimensions of the summary dataset.

```
#12 split-apply-combine summary
library(dplyr)
library(tidyr)

EPAair_summary <- EPAair_wide %>%
  group_by(Site.Name, Month, Year) %>%
  summarise(
    mean_Ozone = mean(Ozone, na.rm = TRUE),
    mean_PM25 = mean(`PM2.5`, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  drop_na(mean_Ozone) # keep rows only where ozone mean exists; pm2.5 may be NA

#13
dim(EPAair_summary)
```

```
## [1] 239    5
```

14. Why did we use the function `drop_na` rather than `na.omit`? Hint: replace `drop_na` with `na.omit` in part 12 and observe what happens with the dimensions of the summary data frame.

Answer: When I use `drop_na`, the function returned more rows than using `na.omit`. The instruction indicated that it's ok to have missing mean PM2.5 values in the result, but using `na.omit` also

eliminate the missing PM2.5 value, but using `drop.na`, we can only eliminate the missing mean Ozone value. In conclusion, `drop_na(mean_Ozone)` is useful here because this only removes where `mean_Ozone` is NA, but `na.omit()` removes all the NA.